



ResNet


Going deeper and deeper...



ResNet



Basic block



Bottleneck block


He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. In Proceedings of the IEEE conference on computer vision and pattern recognition (pp. 770-778).

Practice

- Run "7.4. Build ResNet from scratch.ipynb"



Basic loop



```
class BasicBlock(nn.Module):
    expansion = 1
    def __init__(self, inplanes, planes, stride=1, downsample=None):
        super(BasicBlock, self).__init__()
        self.conv1=conv3x3(inplanes,planes,stride)
        self.bn1=nn.BatchNorm2d(planes)
        self.relu=nn.ReLU(inplace=True)
        self.conv2=conv3x3(planes,planes)
        self.bn2=nn.BatchNorm2d(planes)
        self.downsample=downsample
        self.stride=stride

        if(stride!=1 or inplanes!=planes*self.expansion):
            self.downsample=nn.Sequential(
                nn.Conv2d(inplanes,planes*self.expansion,kernel_size=1,stride
                nn.BatchNorm2d(planes*self.expansion),
            )

    def forward(self, x):
        residual = x
        out = self.conv1(x)
        out = self.bn1(out)
        out = self.relu(out)
        out = self.conv2(out)
        out = self.bn2(out)

        # Downsample:feature Map size/2 || Channel increase
        if (self.downsample is not None):
            residual = self.downsample(x)
        print("out= ", out.shape, "residual= ", residual.shape)
        out+=residual
        out=self.relu(out)
        return out
```

Batch Normalization


Applies Batch Normalization over a 4D input (a mini-batch of 2D inputs with additional channel dimension) as described in the paper [Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift](#).

$$y = \frac{x - \mathbb{E}[x]}{\sqrt{\text{Var}[x] + \epsilon}} * \gamma + \beta$$

- The mean and standard-deviation are calculated per-dimension over the mini-batches.
- By default, the elements of γ are set to 1 and the elements of β are set to 0.

<https://pytorch.org/docs/stable/generated/torch.nn.BatchNorm2d.html>

Batch Normalization



<https://medium.com/ching-i/batch-normalization-%E4%BB%8B%E7%B4%B9-135a24928f12>

```

class MyResNet(nn.Module):
    def __init__(self, block, layers, num_classes=2):
        super(MyResNet, self).__init__()
        self.inplanes = 64
        self.dilation = 1
        self.conv1=nn.Conv2d(3,self.inplanes,kernel_size=7,stride=2,
        self.maxpool=nn.MaxPool2d(kernel_size=3,stride=2, padding=1)
        self.layer1=self._make_layer(block,64,layers[0])
        self.layer2=self._make_layer(block,128,layers[1],stride=2)
        self.avgpool=nn.AdaptiveAvgPool2d((1,1))
        self.fc=nn.Linear(128*block.expansion,num_classes)
        self.linear=nn.Linear(128*block.expansion,num_classes)

    def _make_layer(self, block, planes, blocks, stride=1):
        layers=[]
        layers.append(block(self.inplanes,planes,stride))
        self.inplanes=planes*block.expansion

        for i in range(1,blocks):
            layers.append(block(self.inplanes,planes))
        return nn.Sequential(*layers)

    def forward(self, x):
        x=self.conv1(x)
        x=self.maxpool(x)
        x=self.layer1(x)
        x=self.layer2(x)
        x=self.avgpool(x)
        x=torch.flatten(x, 1)
        x=self.fc(x)
        return x

```

```

model=MyResNet(BasicBlock,[1,1]).to(device)
print(model)

```

```

MyResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (layer2): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
  (fc): Linear(in_features=128, out_features=2, bias=True)
  (linear): Linear(in_features=128, out_features=2, bias=True)
)

```


Practice – Draw the structure of MyResNet

```
MyResNet(  
    (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)  
    (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
```

```
[14]: out1=model.conv1(imageTensor.to(device))  
      print(out1.shape)
```

```
torch.Size([1, 64, 112, 112])
```

```
[15]: out2=model.maxpool(out1)  
      print(out2.shape)
```

```
torch.Size([1, 64, 56, 56])
```

Practice – Draw the structure of MyResNet

```
(layer1): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
```

```
[16]: out3=model.layer1(out2)

out= torch.Size([1, 64, 56, 56]) residual= torch.Size([1, 64, 56, 56])
```

Practice – Draw the structure of MyResNet

```
(layer2): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
)
```

```
[17]: out4 = model.layer2(out3)

      out= torch.Size([1, 128, 28, 28]) residual= torch.Size([1, 128, 28, 28])
```

Practice – Draw the structure of MyResNet

```
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))  
(fc): Linear(in_features=128, out_features=2, bias=True)  
(linear): Linear(in_features=128, out_features=2, bias=True)
```

```
[18]: out5= model.avgpool(out4)  
      print(out5.shape)  
  
      torch.Size([1, 128, 1, 1])
```

```
[19]: out6=torch.flatten(out5,1)  
      print(out6.shape)  
  
      torch.Size([1, 128])
```

```
[20]: out7 = model.fc(out6)  
      print(out7)  
  
      tensor([[ -0.0661, -0.1440]], device:
```

Practice – Load pre-trained ResNet

In [2]: `import torchvision`

`model = torchvision.models.resnet18(pretrained=True)`

Downloading: "<https://download.pytorch.org/models/resnet18-5c106cde.pth>" to

HBox(children=(FloatProgress(value=0.0, max=46827520.0), HTML(value='')))

ResNet


```
ResNet(  
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)  
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  (relu): ReLU(inplace=True)  
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)  
  (layer1): Sequential(  
    (0): BasicBlock(  
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
      (relu): ReLU(inplace=True)  
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    )  
    (1): BasicBlock(  
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
      (relu): ReLU(inplace=True)  
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    )  
  )  
)
```

ResNet

```
(layer2): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
```

Why deep ?

With same number of parameters, which NN is better?



Deep is better

Layer X Size	Word Error Rate (%)	Layer X Size	Word Error Rate (%)
1 X 2k	24.2		
2 X 2k	20.4		
3 X 2k	18.4		
4 X 2k	17.8		
5 X 2k	17.2	1 X 3772	22.5
7 X 2k	17.1	1 X 4634	22.6
		1 X 16k	22.1

Why?

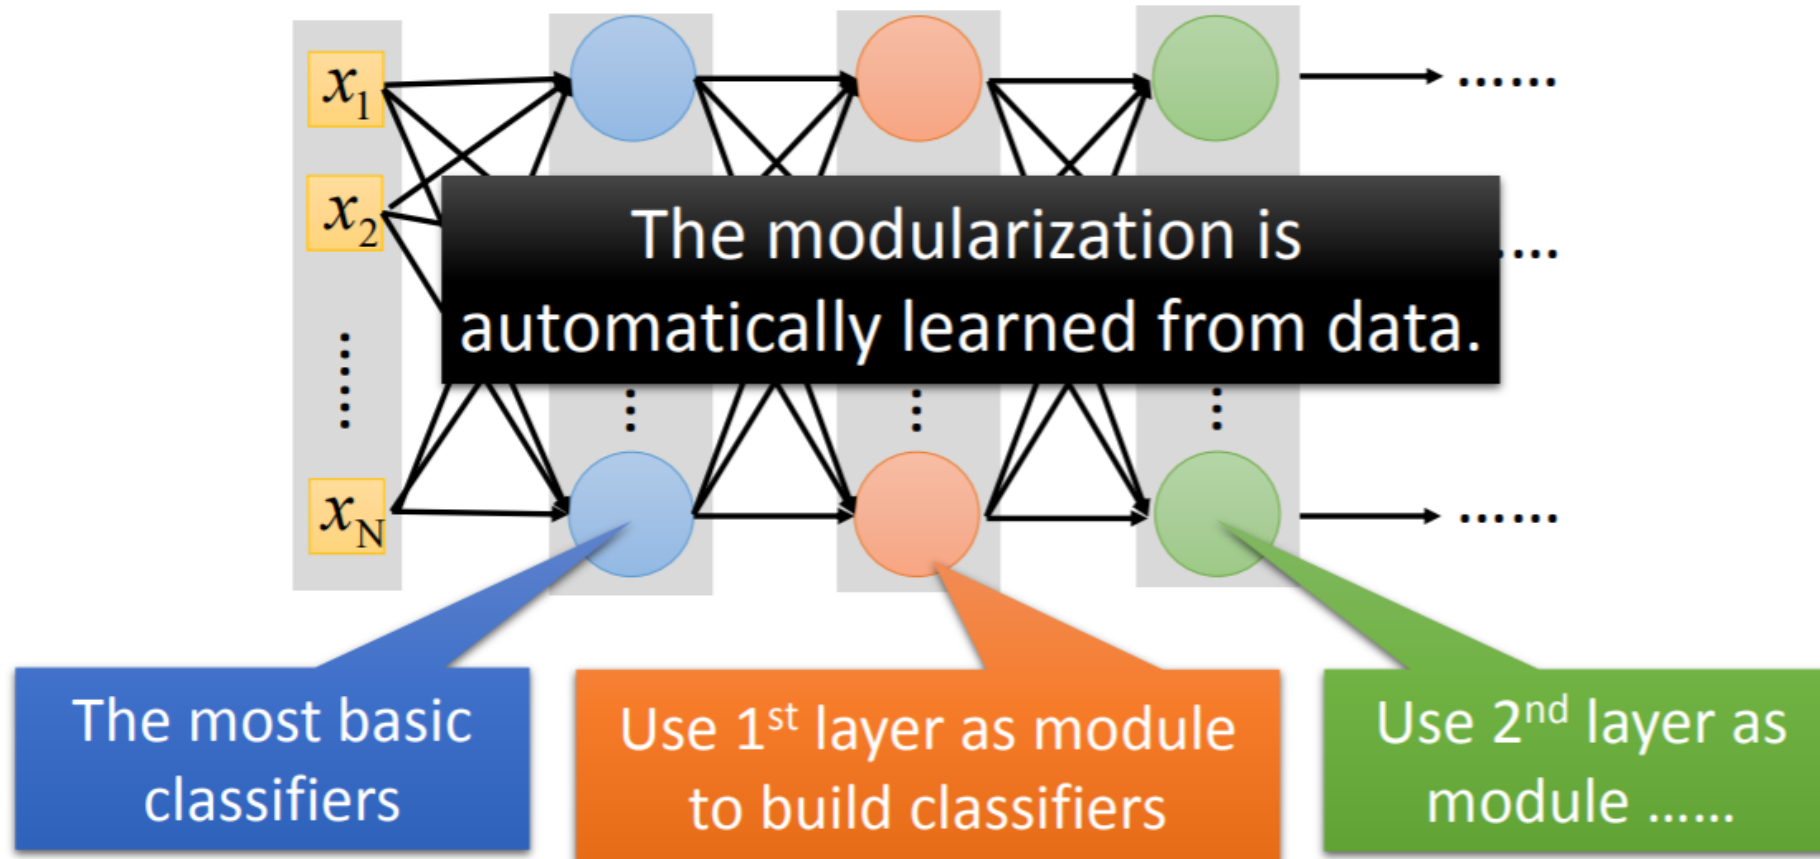
deep + thin

short + fat

Seide, Frank, Gang Li, and Dong Yu. "Conversational Speech Transcription Using Context-Dependent Deep Neural Networks." *Interspeech*. 2011.

Reason 1 – Modularization

- Deep → Modularization → Less training data?




Universality theorem

Any continuous function f

$$f : R^N \rightarrow R^M$$

Can be realized by a network
with one hidden layer

(given **enough** hidden neurons)



Reference for the reason:

<http://neuralnetworksanddeeplearning.com/chap4.html>

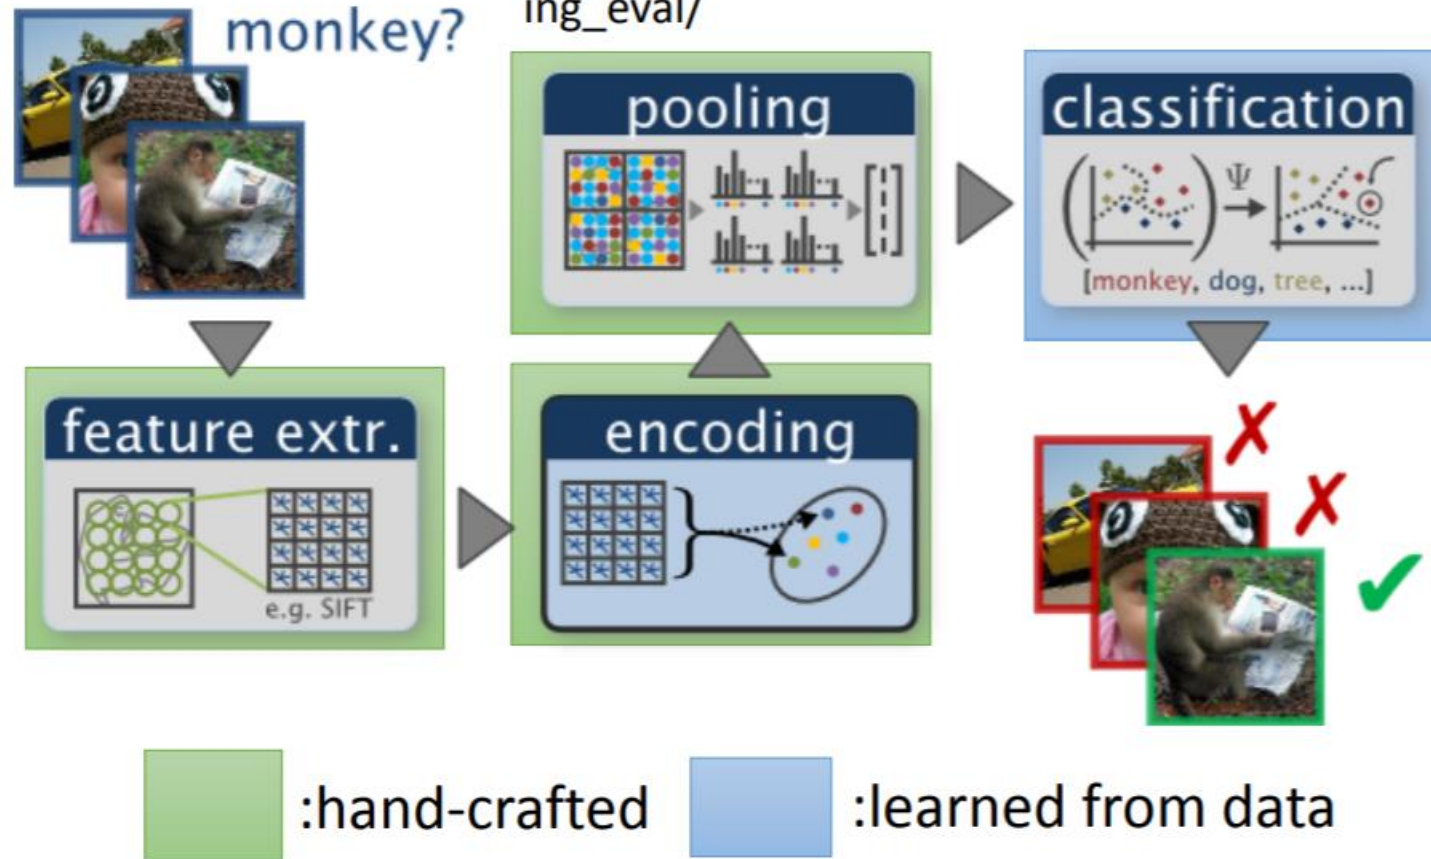
Yes, shallow network can represent any function.

However, using deep structure is more effective.

Reason 2: End-to-end learning


- Shallow Approach

http://www.robots.ox.ac.uk/~vgg/research/encoding_eval/



End-to-end learning

- Deep Learning



Reason 3 - Easier to handle complex task

- Very similar input, different output




- Very different input, similar output



Easier to handle complex task with DL

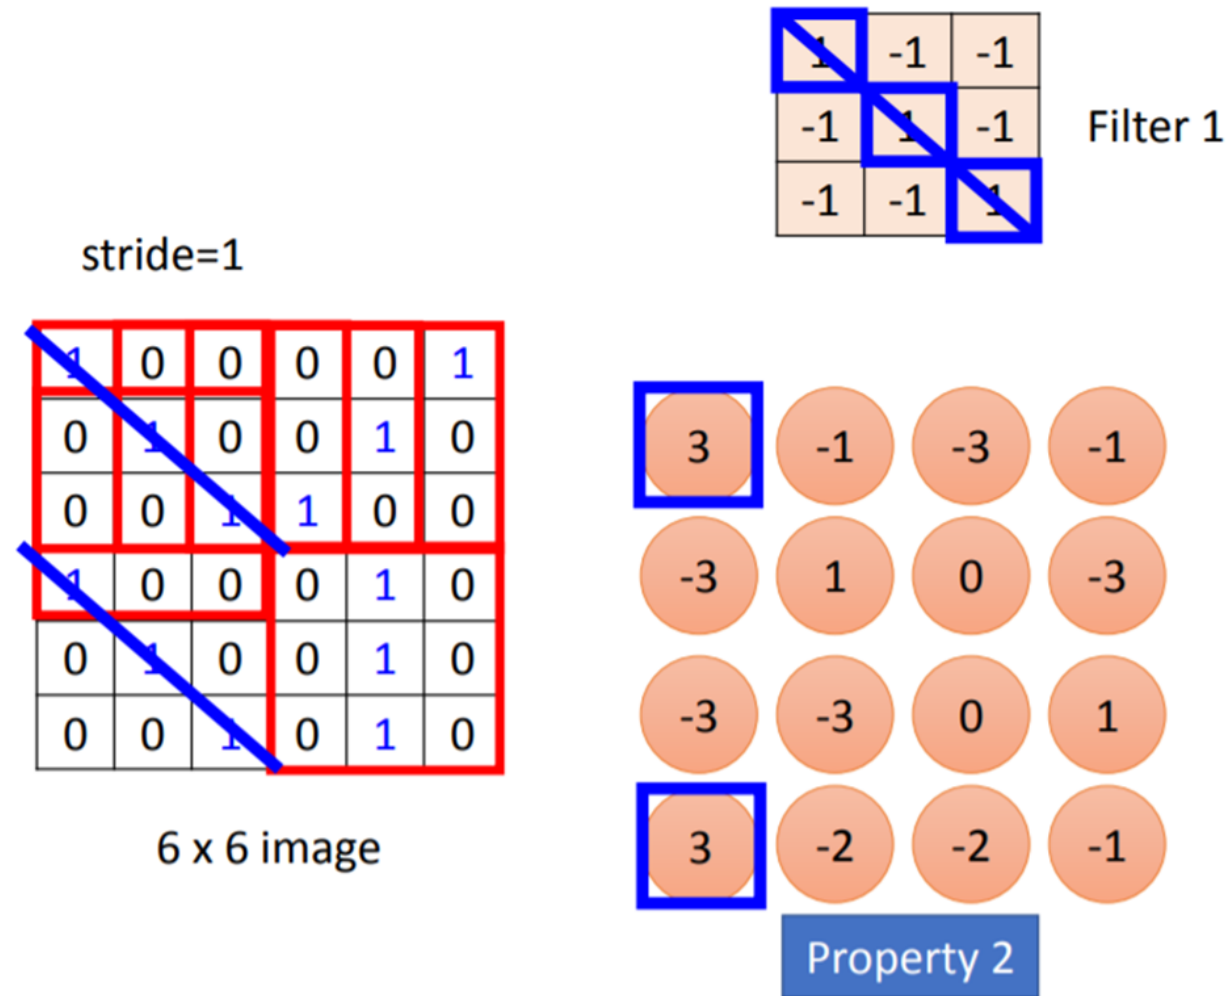
MNIST



How to implement
this in PyTorch?

What does CNN learn?

Recap – Filter searches a particular pattern in different regions and summarize the results in feature map





Only the weight of the 1st convolution filters can be directly visualized.
How to interpret the filter weights of other convolution layers?

How to use
gradient ascent to
implement this in
PyTorch?

The output of the k-th filter is a
11 x 11 matrix.

Degree of the activation
of the k-th filter: $a^k = \sum_{i=1}^{11} \sum_{j=1}^{11} a_{ij}^k$

$$x^* = \arg \max_x a^k \text{ (gradient ascent)}$$




With MNIST data set, in the convolution layer, the filters detects a particular texture pattern.

The output of the k-th filter is a 11 x 11 matrix.

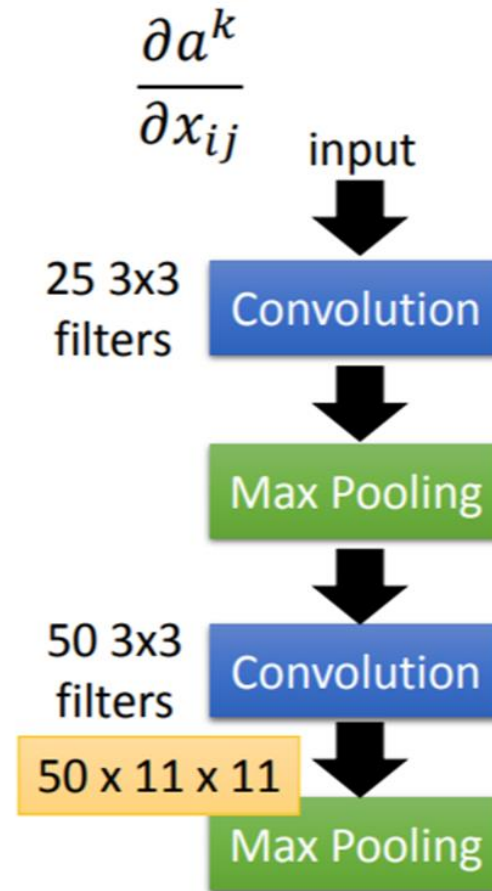
Degree of the activation of the k-th filter:

$$a^k = \sum_{i=1}^{11} \sum_{j=1}^{11} a_{ij}^k$$

$x^* = \arg \max_x a^k$ (gradient ascent)




How to
implement this in
PyTorch?




In the hidden layer of the fully-connected NN, each neuron detects an overall pattern in the picture rather than a particular texture pattern.

Find an image maximizing the output of neuron:

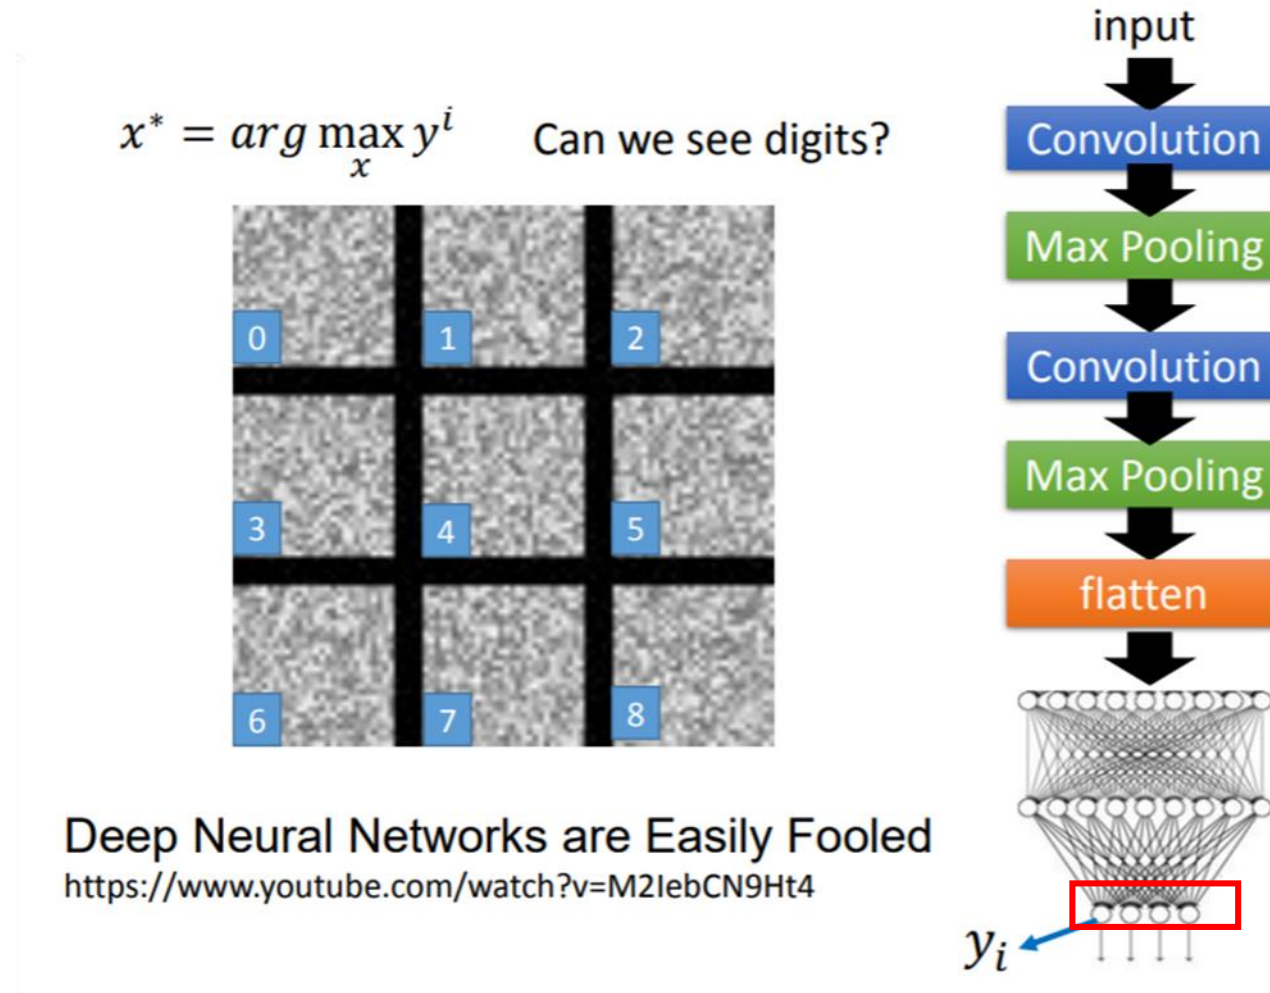
$$x^* = \arg \max_x a_j$$



Each figure corresponds to a neuron



If we watch the output layer node, it is easy to see that CNN is easily fooled.




HOW TO CONFUSE MACHINE LEARNING



Adding regularization to the objective function to force most pixels be "NO INK"


$$x^* = \arg \max_x y^i$$



Here white pixels indicate ink, and black pixels indicate "NO INK".

$$x^* = \arg \max_x \left(y^i - \sum_{i,j} |x_{ij}| \right)$$

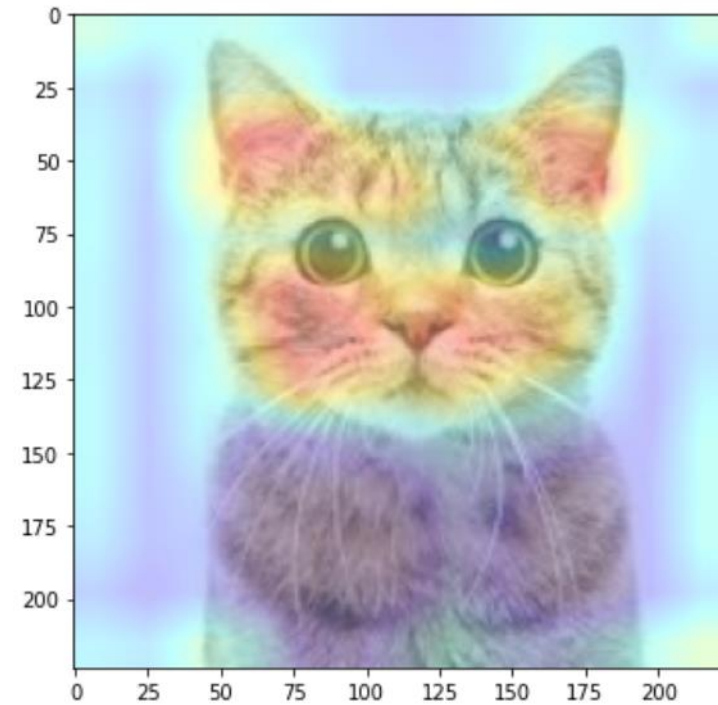

Over all pixel values



L1 regularization to force $x_{ij}=0$, i.e., force most pixels to be black, NO INK (as only small part of the image has ink)

Practice – What does CNN learn?

- Run "6.5 GradCAM.ipynb"



HW5 (3)


	Class index predicted by the model	Class index you assigned
AlexNet		
VGG		
ResNet18		





Use CNN in Alpha GO

- Some patterns are much smaller than the whole image

Alpha Go uses 5 x 5 for first layer



- The same patterns appear in different regions.



Use CNN in Alpha GO

Neural network architecture. The input to the policy network is a $19 \times 19 \times 48$ image stack consisting of 48 feature planes. The first hidden layer zero pads the input into a 23×23 image, then convolves k filters of kernel size 5×5 with stride 1 with the input image and applies a rectifier nonlinearity. Each of the subsequent hidden layers 2 to 12 zero pads the respective previous hidden layer into a 21×21 image, then convolves k filters of kernel size 3×3 with stride 1, again followed by a rectifier nonlinearity. The final layer convolves 1 filter of kernel size 1×1 with stride 1, with a different bias for each position, and applies a softmax function. The **Alpha Go does not use Max Pooling** Extended Data Table 3 additionally show the results of training with $k = 128, 256$ and 384 filters.